

explains several key concepts related to machine learning and computer vision, particularly focusing on how to classify handwritten digits using TF.Learn.

Here's an explanation of the core concepts:

- **Problem Statement: Multi-class Classification** The central problem addressed is **classifying handwritten digits** from the **MNIST dataset**, which is described as the "Hello World of computer vision". This is a **multi-class classification problem**, meaning the goal is to predict which of several possible digits (0-9) an image represents.

- **The MNIST Dataset** The MNIST dataset is a collection of **thousands of labelled images of handwritten digits**. It is pre-divided into:

- A **training set** with 55,000 images.
- A **test set** with 10,000 images.
- The images are **low resolution, specifically 28 by 28 pixels in grayscale**.
- They are also **properly segmented**, meaning each image contains exactly one digit. Some examples are clearly drawn, while others show a variety of handwriting samples that are harder to recognise.

- **Features for Image Classification** When working with images, the classifier uses **raw pixels as features**. This is because extracting more complex features like textures and shapes from images is difficult.

- A 28 by 28-pixel image contains **784 pixels**, which translates to **784 features**.
- These images are used in a **flattened representation**, meaning a 2D array of pixels is converted into a 1D array by unstacking the rows and lining them up. This reshaping is necessary for the classifier but needs to be reversed to display the image.

- **The Linear Classifier (How it Works)** The tutorial demonstrates using a **linear classifier** from TF.Learn.

- **Parameters:** The classifier requires two main parameters: the number of **classes (10, one for each digit)** and information about the **features (784 pixels)**.

- **High-Level Overview:** You can think of the classifier as **adding up evidence** for each type of digit.

- **Input and Output Nodes:**

- There is **one input node for each feature (pixel)** in the image, so 784 input nodes.

- There is **one output node for each digit** the image could represent, so 10 output nodes.
- **Weights and Connections:** The input and output nodes are **fully connected**, and **each connection has a weight**.
- **Classification Process:** When classifying an image, each pixel's intensity flows into its input node, travels along the edges, and is **multiplied by the weight on that edge**. The output nodes then **gather evidence** that the image represents a particular digit.
- **Evidence Calculation:** The evidence for an output node is calculated by **summing the value of the pixel intensities multiplied by their respective weights**.
- **Prediction:** The image is predicted to belong to the output node with the **most evidence**.
- **Importance of Weights:** The **weights are crucial**; by setting them properly, accurate classifications can be achieved.
- **Training (Fitting the Model):** The process begins with **random weights**, which are then **gradually adjusted towards better values**. This adjustment happens within the fit method of the classifier.
- **Evaluating and Making Predictions** Once trained, the model can be evaluated using the evaluate method. In the tutorial, the classifier correctly classifies **about 90% of the test set**. The trained model can also be used to **make predictions on individual images**.
- **Visualising the Learned Weights** The tutorial demonstrates how to **visualise the weights** that the classifier learns.
 - **Representation:** Positive weights are typically drawn in red, and negative weights in blue.
 - **Interpretation:** The weights provide insight into how the classifier "sees" the digits. For instance, a pixel that is almost always filled in for a 'one' (like a middle pixel) would have a **high positive weight (red)** for the 'one' output, indicating strong evidence. Conversely, a pixel that is empty for a 'zero' but might be filled for other digits (like a middle pixel) would have a **negative weight (blue)** for the 'zero' output, indicating evidence against it being a zero if that pixel is filled.
 - By looking at the visualised weights for each class, you can **almost see outlines of the digits** drawn in red.
 - The visualisation is possible because the classifier learns 10 weights for each of the 784 pixels (one for each digit), which are then reshaped back into a 2D array.
- **Environment Setup** The tutorial also briefly covers setting up the environment using **Docker to install TensorFlow**. This involves opening the Docker Quickstart terminal, noting an IP address,

launching a Docker container with a TensorFlow image from Docker Hub, and then accessing an **IPython notebook** in a browser via the noted IP address and port 8888. The necessary libraries, including matplotlib for displaying images and TF.Learn for training, are pre-installed with the Docker image.

The concepts laid out in this tutorial provide a foundational understanding of building and understanding a simple image classifier using a linear model, and it sets the stage for more complex methods like deep learning